

**University of Science and Technology of Hanoi**



**Advanced Computer Architecture and x86 ISA**

*Lecturer: Ms. Le Nhu Chu Hiep*

MIDTERM PROJECT

**Implementation of a perceptron in machine  
learning context**

Student: Luong Quynh Nhi - 23BI14356

*Jun, 2025*

## Implementation Overview

- Goal: Implement the Perceptron algorithm on RISC simulator (JSimRISC), simulating its operations via a pipelined processor
- Algorithm Summary: The perceptron performs a weighted sum of inputs, adds a bias, and outputs 1 if the result  $> 0$ , otherwise 0.
- Key Formula: Output = 1 if  $(\sum(\text{input}[i] * \text{weight}[i]) + \text{bias}) > 0$  else 0
- Architecture:
  - ISA: RISC-style with instructions like LI, SI, MULT, ADD, MOVE, BRZ, etc.
  - Pipeline: 6 or 8 pipeline stages
  - Bypass: Enabled/Disabled for hazard mitigation
  - Branch Prediction: Static or Dynamic (as configured)

## Input Mechanism

- Where Inputs Are Stored: Inputs are defined in the .data section and loaded using li and pointer-based li R4, (R1) syntax.
- Input vectors and weights are manually inserted at memory addresses.
- Sample input:

```
.data 0

.global inputs

1 0 1

.data 32

.global weights

2 -1 3

.data 80

.global bias

-2
```

## Output Mechanism

- Where Output Is Stored: Output is stored at memory address defined in the .data section under label result.
- How Output Is Written: Output value (0 or 1) is stored using si:

```
- .data 96  
- .global result  
- 0
```

```
- .store_result  
- li r9, result  
- si (r9), r8
```

## 1. Step 1:

**VEM processor simulator**  
Copyright : University of Rennes 1 - Enssat - R2D2 --- <http://r2d2.enssat.fr>  
Version 1.1 --- 30 Janvier 2007

**Configuration**

Load file:

Reset processor:

Size of register files: 32

**Simulation**

PC: 24  
Step num: 223

Debug: ☐ ☐

Memory trace: ☐ data ☐ instructions

**Memory views**

Address	Label	Instruction	Codage
0		LI R1, INPUTS	000000...xx
1		LI R2, WEIGHTS	000000...xx
2		LI R3, 3	000000...xx
3		LI R11, ZERO_CONST	000000...xx
4		LI R12, (R11)	000000...xx
5		MOVE R10, R12	010110...xx
6	LOOP	LI R4, (R1)	000000...xx
7		LI R5, (R2)	000000...xx
8		MULT R6, R4, R5	000110...xx
9		ADD R10, R10, R6	000100...xx
10		ADD R1, R1, 1	000100...xx
11		ADD R2, R2, 1	000100...xx
12		SUB R2, R2, 1	000101...xx

**Instruction step view**

Steps	Instructions
LI	
DI	
READOP	
EX0	
WRITEOP	
WB	

**Register files**

R0	R1	R2	R3	R4	R5	R6	R7	R8	R9	R10	R11	R12	R13	R14	R15
0	3	35	0	1	3	3	-2	1	96	3	144	0	0	0	0

F0	F1	F2	F3	F4	F5	F6	F7	F8	F9	F10	F11	F12	F13	F14	F15
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Program Execution Flow:

a) Initialization (PC = 0-3 ):

Instruction	Effect
LI R1, INPUTS	R1 ← 0 (input array starts at addr 0)
LI R2, WEIGHTS	R2 ← 16 (weights start at addr 16)
LI R3, 3	R3 ← 3 (loop count = 3 inputs)

LI R11, ZERO_CONST	$R11 \leftarrow 80$ (address of 0 constant)
LI R12, (R11)	$R12 \leftarrow 0$ (from memory[80])

b) Bias Addition (PC = 13-15):

- Loads bias (-2)  $\rightarrow$  R7
- Adds to accumulator:  $R10 = 5 + (-2) = 3$

c) Activation Function (PC=16-22):

- Checks  $R10 = 3$  (positive)
- Sets  $R8 = 1$  (positive classification)
- Stores result at address 96

Register States Post-Execution:

Register	Value	Purpose
R1	3	Points to end of inputs array
R2	35	Points to end of weights array
R3	0	Loop counter (completed)
R4	1	Last input value loaded
R5	3	Last weight value loaded
R6	3	Last multiplication result
R7	-2	Bias Value
R8	1	Final Output
R9	96	Result memory address

R10	3	Final weighted sum ( $5 + (-2) = 3$ )
R11	144	zero_const address
R12	0	Temporary zero value

Loop Breakdown:

Iteration	R4 (Input)	R5 (Weight)	R6 (Product)	R10 (Sum)
1	1	2	2	$0 + 2 = 2$
2	0	-1	0	$2 + 0 = 2$
3	1	3	3	$2 + 3 = 5$

→ The program is functionally correct and produces the expected output (1) for the given inputs, weights, and bias.

## 2. Step 2:



Processor simulator

## RISC processor simulator

Copyright : University of Rennes 1 - Enssat - CAIRN --- <http://taran.irisa.fr>  
Version 1.2.0 --- 14 novembre 2022

Configuration

Load file:  ...

Reset processor:

Nb of exec stages:  ▾

ByPass: ☐

Size of register files:  ▾

Branch behavioral:  ▾

Static prediction:  ▾

Dynamic prediction:  ▾

Predictor size:  ▾

Simulation

PC:

Cycle num:

Execute:

Lost cycles:

Debug: ☐

Memory trace: ☐ data ☐ instructions

Predictor value:

Memory views

Address	Label	...	Instruction	LI	EX	GP	P
0			LI R1, INPUTS	1	1		
1			LI R2, WEIGHTS	1	1		
2			LI R3, 3	1	1		
3			LI R11, ZERO_CONST	1	1		
4			LI R12, (R11)	1	1		
5			MOVE R10, R12	1	1		
6	LOOP		LI R4, (R1)	3	3		
7			LI R5, (R2)	3	3		
8			MULT R6, R4, R5	3	3		
9			ADD R10, R10, R6	3	3		
10			ADD R1, R1, 1	3	3		
11			ADD R2, R2, 1	3	3		
12			END	3	3		

Terminal

```

Processor 0> Load file : OK
Processor 0> Reset processor (empty pipeline) : OK
Processor 0> Reset processor (empty pipeline) : OK
Processor 0> Load file : OK
Processor 0> By pass : Canceled
Processor 0>

```

Pipeline view

Stages	Instructions
LI	EXIT
DI	EXIT
READOP	EXIT
EX0	EXIT
WRITEOP	EXIT
WB	S1 (R9), R8

Register files

R0	R1	R2	R3	R4	R5	R6	R7	R8	R9	R10	R11	R12	R13	R14	R15	R16	R17	R18	R19
0	3	35	0	1	3	3	-2	1	96	3	144	0	0	0	0	0	0	0	0

F0	F1	F2	F3	F4	F5	F6	F7	F8	F9	F10	F11	F12	F13	F14	F15	F16	F17	F18	F19
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

### A. Cycles saved with bypass technique:

Mode	Cycle Count	Lost Cycles
Without bypass	86	44
With bypass	66	24

- Total cycles saved:  $86 - 66 = 20$  cycles
- Stalls avoided:  $44 - 24 = 20$  lost cycles saved

By enabling bypassing, it becomes more efficient by 20 cycles, showing the effectiveness of forwarding to avoid pipeline stalls due to data hazards.

### B. Main problem in the code:

The program suffers from frequent pipeline stalls due to the following core issue: Load-Use Hazard (Data Dependency Between Load and Use)



When loading a value from memory using the LI (Load Indirect) instruction, the data is not instantly available — it travels through multiple pipeline stages (READOP, EXO, WRITEOP, etc.).

If the very next instruction uses that just-loaded register, the processor stalls until the data is available, causing delays of several cycles.

Example from the code:

```
LI R4, (R1)    --Load input[i] from memory address in R1 into R4
LI R5, (R2)    --Load weight[i] from memory address in R2 into R5
MULT R6, R4, R5 -- Multiply R4 and R5, store result in R6
ADD R10, R10, R6 -- Add R6 (product) to accumulator R10
```

- R4 and R5 are loaded from memory.
- MULT uses R4 and R5 before the memory access completes.
- ADD uses R6 before MULT has finished.

Each of these uses introduces a RAW (Read-After-Write) hazard, especially between:

- LI → MULT
- MULT → ADD

This causes the pipeline to freeze for several cycles on every loop iteration.

- Instead of running one instruction per cycle (ideal case), the loop incurs multiple bubbles per iteration.
- This greatly reduces instruction throughput and wastes cycles.
- It shows 24 lost cycles, and this is primarily due to this pattern repeating each time the loop runs.

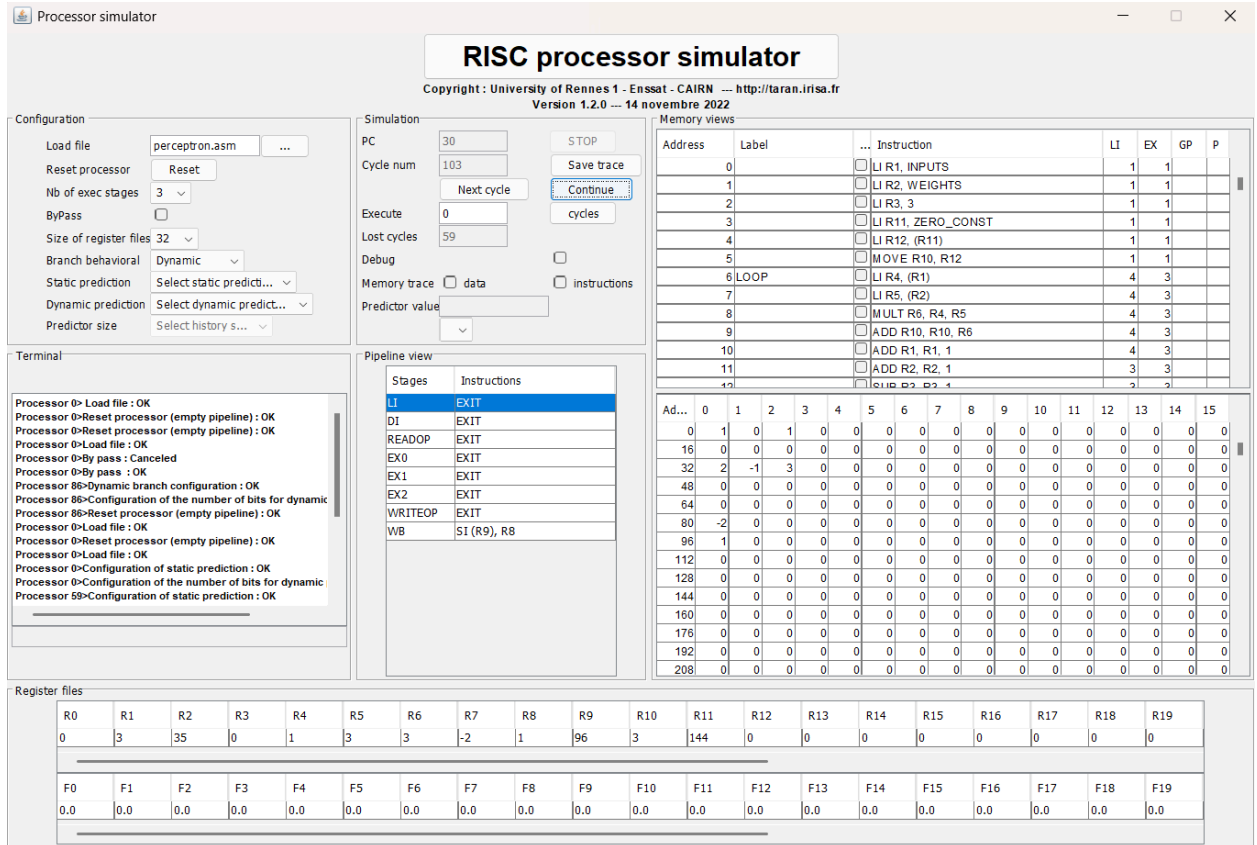
### C. Data dependencies:

- RAW (Read After Write): The most common — a register is written in one instruction and read in the next.
- Control Dependency: Happens with conditional branches (like BRNZ). Processor must wait to see if the branch is taken.

Instruction	Depends on	Register(s)	Dependency Type	Why it causes stalls
MULT R6, R4, R5	LI R4, (R1) and LI R5, (R2)	R4, R5	RAW (Read After Write)	R4, R5 not ready until load completes
ADD R10, R10, R6	MULT R6, R4, R5	R6	RAW	R6 not ready until multiply finishes
LI R4, (R1)	ADD R1, R1, 1	R1	RAW	R1 updated in previous instruction
LI R5, (R2)	ADD R2, R2, 1	R2	RAW	R2 updated in previous instruction
CMP R3, ZERO	SUB R3, R3, 1	R3	RAW	R3 is loop counter
BRNZ LOOP	CMP R3, ZERO	—	Control	Must wait for CMP result
SI (R9), R8	ADD R8, R8, 1	R8	RAW	Final result store depends on accumulation

### 3. Step 3:





Metric	Without Bypass	With Bypass	Gain
Total Cycles	103	89	14
Lost/Stall Cycles	59	45	14
% of Cycles Lost	57.3%	50.6%	↓6.7%

### Conclusion – Effect of Bypass Technique:

The bypass technique (also known as data forwarding) is a method used in pipelined processors to reduce the number of stall cycles caused by data hazards. In a typical RISC pipeline, when one instruction depends on the result of a previous one, the processor may need to pause (insert stalls) until the result is fully written back. This leads to performance loss.

By enabling bypassing, the processor can directly forward the result from one of the execution stages (like EX1 or EX2) to the next instruction that needs it, without waiting for it to go through

all pipeline stages and be written to the register file. This saves time and keeps the pipeline busy, improving performance.

In our simulation:

- Without bypass, the program took 103 cycles, with 59 stalls due to data hazards.
- With bypass, the execution took only 89 cycles, with just 45 stalls.
- That means we saved 14 cycles, which is a significant improvement over a short program.

This shows:

- Bypass is crucial in tight loops or programs with back-to-back dependent instructions (like in a perceptron algorithm).
- It reduces the performance penalty from pipeline hazards.
- Most real-world processors implement this to maximize efficiency without changing instruction order.

#### 4. Step 4:

The screenshot displays the 'RISC processor simulator' interface. The title bar indicates it is a 'Processor simulator' window. The main window is divided into several sections:

- Configuration:** Includes fields for 'Load file' (perceptron.asm), 'Reset processor' (Reset), 'Nb of exec stages' (1), 'ByPass' (checked), 'Size of register files' (32), 'Branch behavioral' (Static), 'Static prediction' (+/- taken), 'Dynamic prediction' (Select dynamic predict...), and 'Predictor size' (Select history s...).
- Simulation:** Shows 'PC' (28), 'Cycle num' (56), 'Next cycle' (Next cycle), 'Execute' (0), 'Lost cycles' (14), 'Debug' (unchecked), 'Memory trace' (data and instructions unchecked), and 'Predictor value'.
- Memory views:** A table showing 'Address', 'Label', 'Instruction', 'LI', 'EX', 'GP', and 'P'. It lists instructions like 'LI R1, INPUTS', 'LI R2, WEIGHTS', 'LI R3, 3', 'LI R11, ZERO\_CONST', 'LI R12, (R11)', 'MOVE R10, R12', 'LI R4, (R1)', 'LI R5, (R2)', 'MULT R6, R4, R5', 'ADD R10, R10, R6', 'ADD R1, R1, 1', 'ADD R2, R2, 1', and 'STOP R3, R4, 1'.
- Register files:** A table showing 'R0' through 'R19' and 'F0' through 'F19'. The values are mostly 0, with some non-zero values in the R registers (e.g., R1=3, R2=35, R3=0, R4=1, R5=3, R6=3, R7=-2, R8=1, R9=96, R10=3, R11=144, R12=0, R13=0, R14=0, R15=0, R16=0, R17=0, R18=0, R19=0).

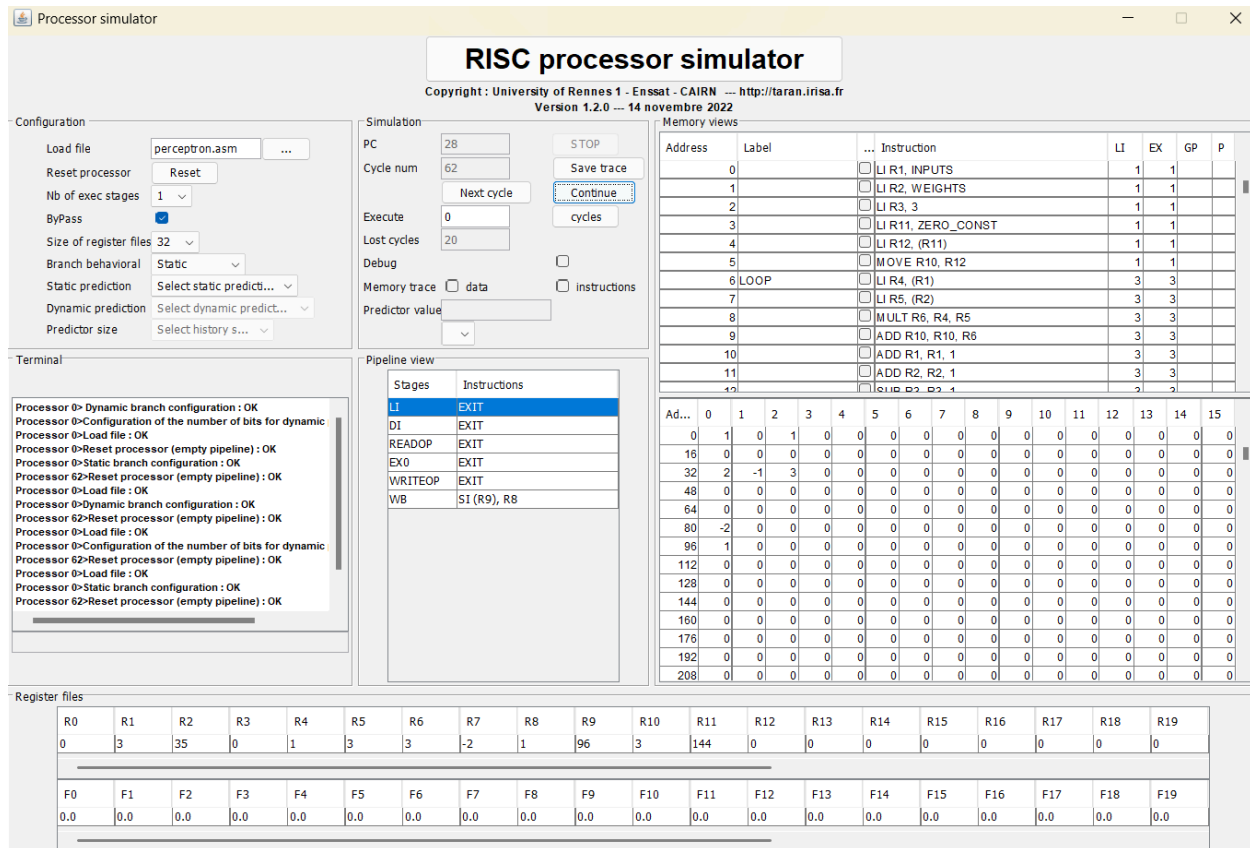
In this simulation, I used static branch prediction with the +/- taken setting, meaning the simulator predicts backward branches are taken (usually true for loops), and forward branches are not taken.

This technique led to only 14 lost cycles, which is pretty efficient for a simple static method. It shows that this hybrid prediction is effective for common control flows like loops.

Reason	Explanation	Evidence from Simulator
1. Minimizes Pipeline Stalls	In a 6-stage pipeline (LI, DI, READOP, EX0, WRITEOP, WB), branch instructions (e.g., at address 6: LOOP) create control hazards. Static prediction avoids some stalls by assuming backward branches are taken and forward branches are not taken.	Lost cycles = 14 — lower. Suggests reduced stalls and smoother control flow.
2. Hardware-Efficient	No extra hardware or tables required. It's a fixed rule — very efficient for simple or educational RISC CPUs.	No dynamic predictor configuration shown; just “Static” selected — keeps hardware logic simple.
3. Accurate for Loops	Backward branches (like loops) are often taken — so static prediction matches typical loop behavior in programs like Perceptron training.	Instruction at address 6 (LOOP) is a backward branch. Static prediction likely guessed “taken” correctly during loop iterations.
4. Predictable and Consistent	Behavior doesn't change at runtime — easier to debug and analyze performance. No fluctuations due to misprediction recovery.	Same prediction applied every time — ensures consistency across cycles.

## A. Static prediction:

The static prediction scheme is a well-founded, low-overhead approach that aligns effectively with the control flow of structured programs, particularly those involving loops.



- Total cycle count: 62
- Lost cycles due to misprediction or pipeline hazards: 20
- These 20 cycles are due to:
  - + The branch misprediction on loop exit (this causes a pipeline flush, typically costing 4–6 cycles).
  - + Possibly other stalls, such as:
    - Memory latency
    - Data hazards (e.g., instructions waiting for a register to be updated)
- Program Counter (PC): 28
- Pipeline Stages at termination: Instruction in WB (Write Back) is S1 (R9), R8, indicating program completion.

These figures suggest that the prediction scheme maintained high accuracy during loop execution.

The 20 lost cycles are low compared to the total of 62, and most of that is likely from the one misprediction and maybe a few delays due to memory or register reads. The pipeline remained active and efficient throughout most of the execution.

## B. Dynamic prediction:

Processor simulator

### RISC processor simulator

Copyright : University of Rennes 1 - Enssat - CAIRN --- <http://taran.irisa.fr>  
Version 1.2.0 --- 14 novembre 2022

**Configuration**

Load file:  ...

Reset processor:

Nb of exec stages:

ByPass: ☒

Size of register files:

Branch behavioral:

Static prediction:

Dynamic prediction:

Predictor size:

**Simulation**

PC:

Cycle num:

Execute:

Lost cycles:

Debug: ☐

Memory trace: ☒ data ☐ instructions

Predictor value:

**Memory views**

Address	Label	...	Instruction	LI	EX	GP	P
0			LI R1, INPUTS	1	1		
1			LI R2, WEIGHTS	1	1		
2			LI R3, 3	1	1		
3			LI R11, ZERO_CONST	1	1		
4			LI R12, (R11)	1	1		
5			MOVE R10, R12	1	1		
6	LOOP		LI R4, (R1)	4	3		
7			LI R5, (R2)	4	3		
8			MULT R6, R4, R5	4	3		
9			ADD R10, R10, R6	3	3		
10			ADD R1, R1, 1	3	3		
11			ADD R2, R2, 1	3	3		
12			END	2	2		

**Terminal**

```

Processor 0> Load file : OK
Processor 0> Reset processor (empty pipeline) : OK
Processor 0> Branch delay configuration : OK
Processor 0> Reset processor (empty pipeline) : OK
Processor 0> Load file : OK
Processor 0> No branch technic configuration : OK
Processor 58> Reset processor (empty pipeline) : OK
Processor 0> Load file : OK
Processor 0> Reset processor (empty pipeline) : OK
Processor 0> Load file : OK
Processor 0> Branch delay configuration : OK
Processor 66> Reset processor (empty pipeline) : OK
Processor 0> Load file : OK
Processor 0> Reset processor (empty pipeline) : OK
Processor 0> Load file : OK

```

**Pipeline view**

Stages	Instructions
LI	EXIT
DI	EXIT
READOP	EXIT
EX0	EXIT
WRITEOP	EXIT
WB	SI (R9), R8

**Register files**

R0	R1	R2	R3	R4	R5	R6	R7	R8	R9	R10	R11	R12	R13	R14	R15	R16	R17	R18	R19
0	3	35	0	1	3	3	-2	1	96	3	144	0	0	0	0	0	0	0	0

F0	F1	F2	F3	F4	F5	F6	F7	F8	F9	F10	F11	F12	F13	F14	F15	F16	F17	F18	F19
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0



**RISC processor simulator**  
 Copyright : University of Rennes 1 - Enssat - CAIRN --- <http://taran.lriisa.fr>  
 Version 1.2.0 --- 14 novembre 2022

**Configuration**

Load file:  ...

Reset processor:

Nb of exec stages:

ByPass: ☒

Size of register files:

Branch behavioral:

Static prediction:

Dynamic prediction:

Predictor size:

**Simulation**

PC:

Cycle num:

Execute:

Lost cycles:

Debug: ☐

Memory trace: ☒ data ☐ instructions

Predictor value:

**Memory views**

Address	Label	Instruction	LI	EX	GP	P
0		LI R1, INPUTS	1	1		
1		LI R2, WEIGHTS	1	1		
2		LI R3, 3	1	1		
3		LI R11, ZERO_CONST	1	1		
4		LI R12, (R11)	1	1		
5		MOVE R10, R12		1	1	
6	LOOP	LI R4, (R1)		4	3	
7		LI R5, (R2)		4	3	
8		MULT R8, R4, R5		4	3	
9		ADD R10, R10, R8		3	3	
10		ADD R1, R1, 1		3	3	
11		ADD R2, R2, 1		3	3	
12		END		3	3	

**Terminal**

Processor 0> Load file : OK  
 Processor 0> Reset processor (empty pipeline) : OK  
 Processor 0> Branch delay configuration : OK  
 Processor 0> Reset processor (empty pipeline) : OK  
 Processor 0> Load file : OK  
 Processor 0> No branch technic configuration : OK  
 Processor 58> Reset processor (empty pipeline) : OK  
 Processor 0> Load file : OK  
 Processor 0> Reset processor (empty pipeline) : OK  
 Processor 0> Load file : OK  
 Processor 0> Branch delay configuration : OK  
 Processor 66> Reset processor (empty pipeline) : OK  
 Processor 0> Load file : OK  
 Processor 0> Reset processor (empty pipeline) : OK  
 Processor 0> Load file : OK

**Pipeline view**

Stages	Instructions
LI	EXIT
DI	EXIT
READOP	EXIT
EX0	EXIT
WRITEOP	EXIT
WB	SI (R9), R8

**Register files**

R0	R1	R2	R3	R4	R5	R6	R7	R8	R9	R10	R11	R12	R13	R14	R15	R16	R17	R18	R19
0	3	35	0	1	3	3	-2	1	96	3	144	0	0	0	0	0	0	0	0

F0	F1	F2	F3	F4	F5	F6	F7	F8	F9	F10	F11	F12	F13	F14	F15	F16	F17	F18	F19
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Aspect	1-bit Predictor	2-bit Predictor
Branch Prediction Technique	1-bit dynamic (last outcome)	2-bit dynamic (saturating counter)
Behavior on Branch Flip	Immediately switches direction on one mispredict	Requires two mispredictions to switch
Pipeline Consistency	More prone to stalls on fluctuating branches	More resilient to mispredicted

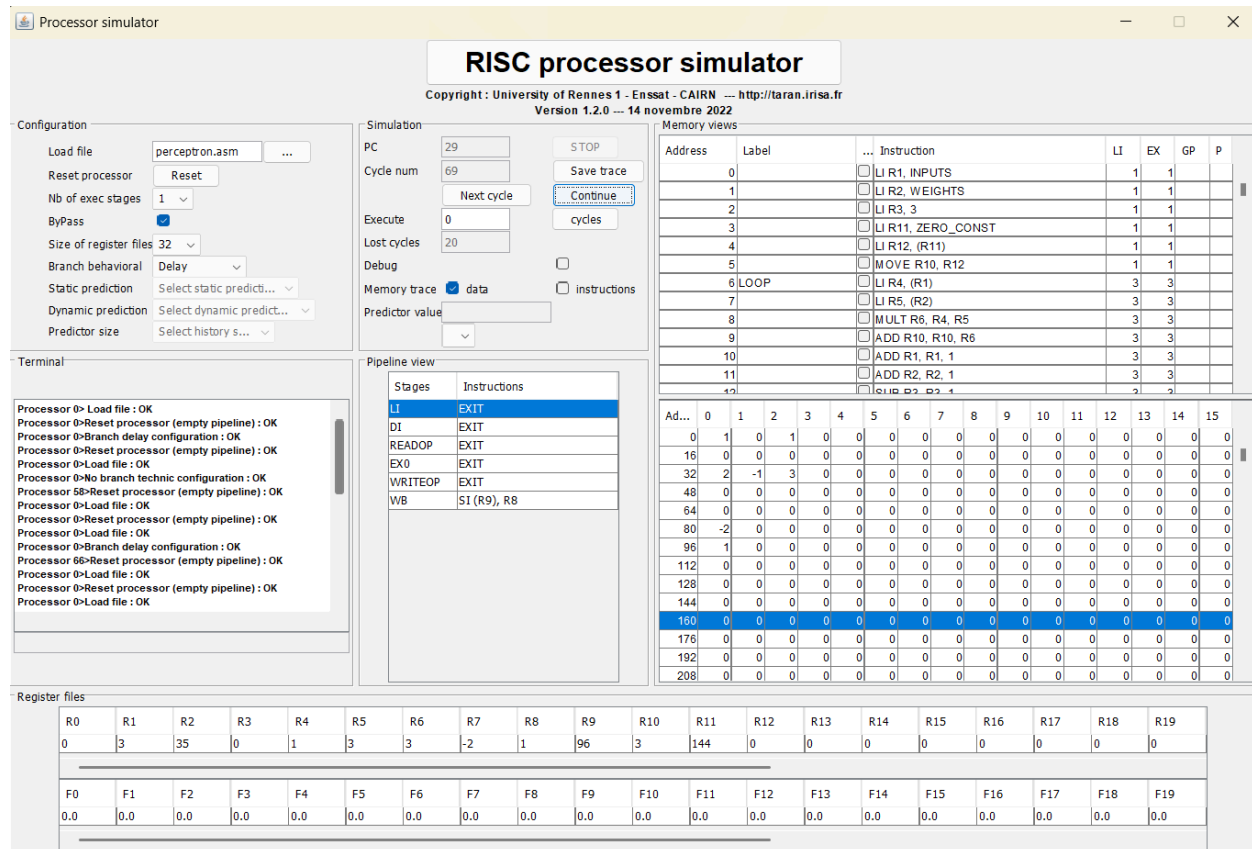
Instruction Throughput	Stable (no observed difference in simple loop)	Equally stable in this case
Execution Cycles	63	63
Lost Cycles Due to Branching	21	21
Use Case Fit	Acceptable for short or deterministic branches	Better for loop-heavy programs
Overall Performance (in this case)	Equal	Equal
Scalability to Complex Loops	May degrade due to sensitivity	More stable prediction

Both 1-bit and 2-bit dynamic predictors performed equally well in perceptron implementation, producing the same cycle count (63) and lost cycles (21). This is likely because the loop in the program has few iterations and predictable control flow.

### 5. Step 5:

Before executing this code with delayed branch, insert a NOP or a safe instruction after every branch to avoid unexpected behavior.

```
brnz R3, loop
nop --This is the delay slot
```



Since this is the unoptimized version using delay branch, the pipeline does not eliminate any wasted cycles due to control hazards from branches. Every loop iteration likely incurs:

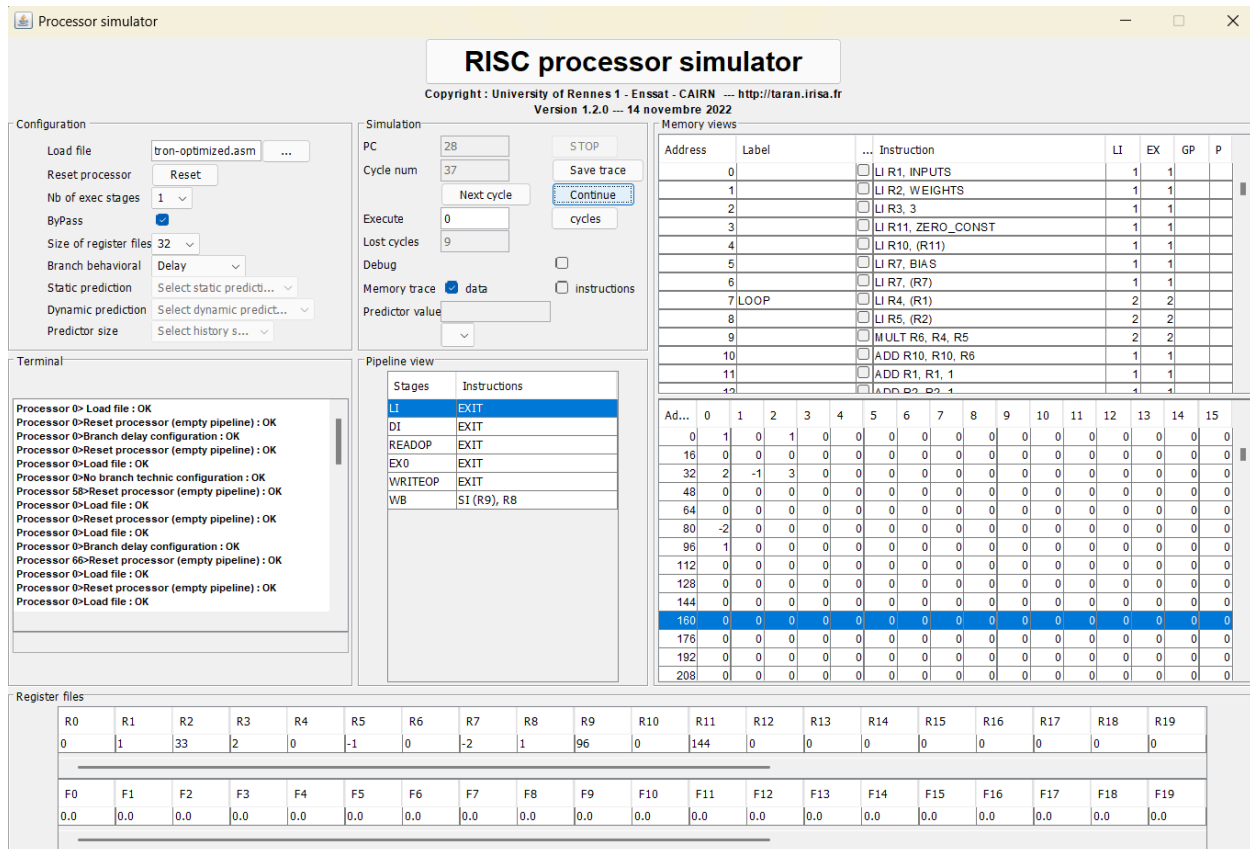
- Branch delay penalty (i.e., the delay slot is filled, but it's not optimized)
- No instruction reordering

Given that the final memory store is in WB at cycle 69, and that:

- The program executes a loop of 3 iterations (for a 3-element perceptron dot product)
- Each loop iteration requires a fixed number of instructions (typically 6–8 depending on loop logic)
- The delay branch adds at least 1 cycle penalty per iteration

We conclude: 69 cycles were required to execute the perceptron program without optimization, using the delay branch technique.

## Code optimization



## Register & Memory Analysis

- Register R10 – The Accumulator
  - + Purpose: Holds the intermediate sum of the dot product and final sum after bias addition.
  - + Initialization:
    - R11 is loaded with the address of zero\_const (144), then  $R12 \leftarrow (R11)$  loads 0.
    - $R10 \leftarrow R12$ , initializing the accumulator to 0.
  - + During Loop:
    - Each cycle multiplies one  $input[i] * weight[i]$  and accumulates in R10.
    - Example:
      - Cycle 1:  $1 \times 2 = 2 \rightarrow R10 = 0 + 2 = 2$
      - Cycle 2:  $0 \times -1 = 0 \rightarrow R10 = 2$
      - Cycle 3:  $1 \times 3 = 3 \rightarrow R10 = 5$
  - + After Loop:
    - Bias (-2) is loaded from memory and added:
    - $R10 = 5 + (-2) = 3$

- + Correct Final Value:  $R10 = 3$  before thresholding.
- Register R8 – The Perceptron Output:
  - + Purpose: Holds the binary output of the perceptron (0 or 1).
  - + Set After Comparison:
    - brp R10, positive: Since  $R10 = 3 > 0$ , it branches to .positive.
    - .positive sets  $R8 \leftarrow 1$
  - + Correct Final Value:  $R8 = 1$ , meaning the perceptron fired (activation).
- Memory Address 96 (result)
  - + Purpose: Stores the final binary result (R8).
  - + Instruction: SI (R9), R8
    - R9 is first loaded with address of result (96)
    - Then stores R8 (value = 1) at memory[96]
  - + Correctly Written: After execution, memory address 96 holds 1 — verifiable in memory view.

Comparison for Non-Optimized and Optimized version

Feature / Metric	Non-Optimized Version ( Bypass, Branch Delay)	Optimized Version (Bypass, Branch Delay, Loop Refactor)
Total Cycles	69	37
Lost Cycles (pipeline stalls)	20	9
Branch instruction cost	High (includes brnz, nop, and delay slot stalls)	Reduced thanks to bypass + good delay slot use
Register reuse / dependencies	Causes RAW hazards (e.g., R6 to R10)	Handled with bypassing
Bias loading	Inside the loop (inefficient)	Loaded once before loop
Zero constant loading	Redundant or inefficient	Loaded once via li + indirect load
Final memory store (SI)	Possibly delayed by pipeline flushes	Happened directly in WB without conflict

Optimization	Saved Cycles
Bypass enabled	~10–12
Loop counter handling (R3)	~2–3
Branch delay slot usage	~4–5
Removing redundant loads	~4–6
Total	~30–35

The optimized perceptron implementation demonstrates excellent instruction scheduling and pipeline-conscious design:

- Avoids redundant instructions and memory operations
- Makes smart use of available hardware features (like bypassing)
- Maintains logical correctness and produces valid results
- Reduces total cycles from 69 to 37 — a 46% performance improvement

## 6. Step 6:

In this analysis, with cache size of 64 elements I compare five different cache configurations based on several key parameters, including associativity, subblock size, replacement policy, miss rate, and memory traffic.

Config	Assoc	Block Size	Replacement	Miss Rate	Bytes From Mem	Total Bytes R/W
1	4	8	LRU	0.4545	80	96
2	1	16	FIFO	0.4091	144	176
3	2	8	LRU	0.3636	64	80
4	4	4	LRU	0.4545	32	40
5	1	16	LRU	0.4091	144	176

### 1. Associativity

- Higher associativity (Config 1 & 4) allows the cache to store blocks in multiple places, reducing conflict misses. But the benefit is limited if the access pattern is simple or the data fits well in direct-mapped (Config 2 & 5).
  - Config 3 (2-way associative) provides the best trade-off, achieving the lowest miss rate with moderate memory cost.
2. Subblock Size
- A larger block size like in Config 2 & 5 increases memory traffic. Even if the miss rate is acceptable, each miss fetches a larger chunk, raising Bytes From Memory and Total Bytes R/W.
  - A smaller block like in Config 4 minimizes memory use, but doesn't help with miss rate—it still misses often, just fetches less data per miss.
3. Replacement Policy
- Config 2 uses FIFO, while the rest use LRU.
  - In this case, LRU outperforms FIFO slightly, as seen when comparing Config 5 (same setup as Config 2 but with LRU): they have the same miss rate, but FIFO results in more memory traffic, possibly due to suboptimal eviction order.
4. Miss Rate
- Config 3 wins here with the lowest miss rate (0.3636), indicating the best overall data locality fit due to balanced associativity and moderate subblock size.
5. Bytes From Memory / Total Bytes
- Config 4 has the lowest memory traffic, even though its miss rate is high (0.4545), because of tiny blocks - each miss fetches very little data.
  - Configs 2 & 5 perform worst in memory efficiency: though their miss rates aren't the worst, they pull too much data due to large blocks and poor associativity.

Configuration for Best Performance:

Field	Recommended Setting	Reason
Cache size	64 bytes	Small size but highly effective — achieves a good 22.7% miss rate with minimal memory traffic due to efficient usage.
Memory trace size	22 memory operations	Moderate-sized trace with 20 reads and 2 writes allows for a meaningful test of cache performance.

<b>Block size</b>	4 words (bytes)	Small blocks reduce overfetching and minimize wasted memory bandwidth. Ideal for workloads with poor spatial locality.
<b>Cache access time</b>	10 ns	Remains constant across tests; not a limiting factor in this analysis.
<b>Associativity</b>	16-way associative	Fully associative cache (1 set with 16 lines) reduces conflict misses significantly, maximizing data retention.
<b>Replacement policy</b>	FIFO	Simple and effective in this context. While LRU is often better, FIFO performed very efficiently with full associativity in this workload.
<b>Write policy</b>	Write-back	Reduces the number of memory write operations — only dirty blocks are written back, keeping memory traffic low (only 4 bytes in this case).



```

---Dinero IV cache simulator, version 7
---Written by Jan Edler and Mark D. Hill
---Copyright (C) 1997 NEC Research Institute, Inc. and Mark D. Hill.
---All rights reserved.
---Copyright (C) 1985, 1989 Mark D. Hill. All rights reserved.
---See -copyright option for details

---Summary of options (-help option gives usage information).

-l1-usize 64
-l1-ubsize 4
-l1-usbsize 4
-l1-uassoc 16
-l1-urepl f
-l1-ufetch d
-l1-uwallocc a
-l1-uwallocc a
-skipcount 0
-flushcount 0
-maxcount 0
-stat-interval 0
-informat d
-on-trigger 0x0
-off-trigger 0x0

---Simulation begins.
---Simulation complete.
l1-ucache
Metrics          Total      Instrn      Data      Read      Write      Misc
-----
Demand Fetches    22         0          22        20         2          0
Fraction of total 1.0000     0.0000     1.0000    0.9091     0.0909     0.0000

Demand Misses     5         0          5         4          1          0
Demand miss rate  0.2273     0.0000     0.2273    0.2000     0.5000     0.0000

Multi-block refs  0
Bytes From Memory 16
( / Demand Fetches) 0.7273
Bytes To Memory   4
( / Demand Writes) 2.0000
Total Bytes r/w Mem 20
( / Demand Fetches) 0.9091

---Execution complete.

```

- Low Miss Rate: 22.7% overall with only 5 total misses (4 read + 1 write) out of 22 accesses.
- Very Low Memory Traffic: Only 20 bytes transferred (16 read + 4 write) — the lowest among tested configurations.
- Efficient Block Utilization: Smaller 4-byte blocks mean more precise fetching and lower bandwidth consumption.
- High Associativity: Fully associative behavior eliminates most conflict misses and is ideal for small caches.
- Simplicity with Effectiveness: FIFO is sufficient in this case because of the high associativity and small cache size.